

## 1. Weather Data Results

1.1) Weather Data Results for each of the versions of sequential and multithreaded program **without** Fibonacci Delay:

No of Threads = 4

SEQ

Average Running Time : 2947.9

Maximum Running Time : 3502

Minimum Running Time : 2853

NO-LOCK

Average Running Time : 1530.8

Maximum Running Time : 1601

Minimum Running Time : 1502

Speed Up : 1.92

COARSE-LOCK

Average Running Time : 1536.0

Maximum Running Time : 1568

Minimum Running Time : 1520

Speed Up : 1.91

FINE-LOCK

Average Running Time : 1557.8

Maximum Running Time : 1613

Minimum Running Time : 1534

Speed Up : 1.89

NO-SHARING

Average Running Time : 1538.2

Maximum Running Time : 1562

Minimum Running Time : 1525

Speed Up : 1.91

1.2) Weather Data Results for each of the versions of sequential and multithreaded program **with** Fibonacci Delay:

No of Threads = 4

SEQ

Average Running Time : 2931.3

Maximum Running Time : 3854

Minimum Running Time : 2770

NO-LOCK

Average Running Time : 1523.5

Maximum Running Time : 1577

Minimum Running Time : 1501

Speed Up : 1.92

COARSE-LOCK

Average Running Time : 1540.8

Maximum Running Time : 1607

Minimum Running Time : 1523

Speed Up : 1.90

FINE-LOCK

Average Running Time : 1562.4

Maximum Running Time : 1688

Minimum Running Time : 1538

Speed Up : 1.87

NO-SHARING

Average Running Time : 1540.6

Maximum Running Time : 1630

Minimum Running Time : 1514

Speed Up : 1.90

### 1.3) Question and Answers

1. Which program version (SEQ, NO-LOCK, COARSE-LOCK, FINE-LOCK, NO-SHARING) would you normally expect to finish fastest and why? Do the experiments confirm your expectation? If not, try to explain the reasons.

**Answer:** I expected the NO-LOCK program version to finish the fastest. It is faster than SEQ because we are processing the data using multiple threads (parallel processing). And it is faster than other multi-threaded versions because we are not blocking the processing while working on shared data or no overhead of merging the results as we have seen in NO-SHARING approach.

Yes, The results confirm my expectation. We got the best Speed up in NO-LOCK program.

2. Which program version (SEQ, NO-LOCK, COARSE-LOCK, FINE-LOCK, NO-SHARING) would you normally expect to finish slowest and why? Do the experiments confirm your expectation? If not, try to explain the reasons.

**Answer:** I expected the SEQ program to finish slowest as we are not doing any parallel processing here. We have a data available in memory and have a processors to work for but due to lack of parallelism we are not using them optimally. We are processing each line one by one here and which is making it slower.

Yes, The results confirm my expectation.

3. Compare the temperature averages returned by each program version. Report if any of them is incorrect or if any of the programs crashed because of concurrent accesses.

**Answer:** I have added a code which compares the averages of each program version and found that NO-LOCK program is giving the incorrect averages and all other approaches are giving the same results. This is because we have not used Locking in this version, due to this multiple threads makes shared data structure inconsistent.

4. Compare the running times of SEQ and COARSE-LOCK. Try to explain why one is slower than the other. (Make sure to consider the results of both B and C—this might support or refute a possible hypothesis.)

**Answer:** COARSE-LOCK program is much faster than SEQ (Speed Up around 1.9) for both B and C. This is because in SEQ approach we are processing each record one by one but in COARSE-LOCK approach we have divided our same task to 4 threads and each thread works on their piece of data.

Though we have a good performance in COARSE-LOCK, if we do very intense work (high time) inside a synchronized block (locked common data structure) then at a time only one thread can work on that part of code which makes it serial and reduces our performance.

5. How does the higher computation cost in part C (additional Fibonacci computation) affect the difference between COARSE-LOCK and FINE-LOCK? Try to explain the reason.

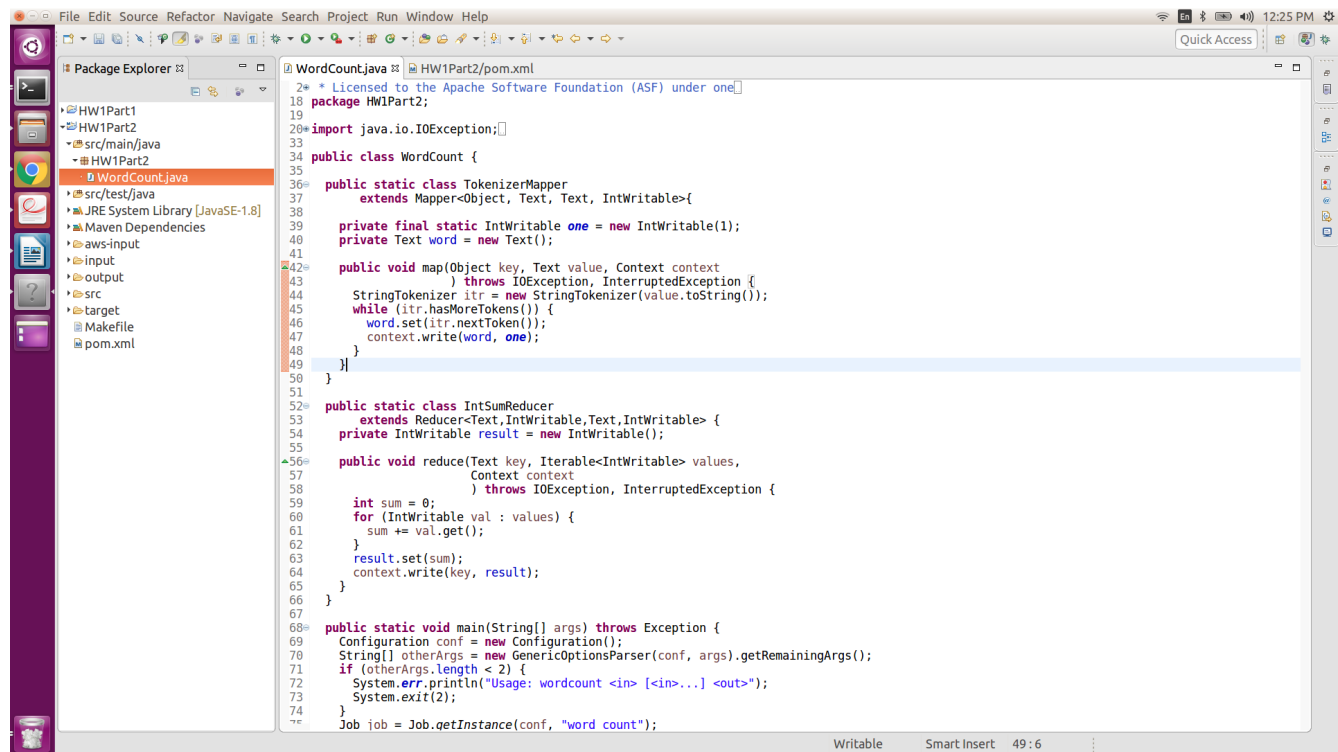
**Answer:** I cannot see much difference in my results but by speed-up measure, I can say adding the higher computation cost in Part C reduced Speed-up of COARSE-LOCK approach from 1.91 to 1.9 while of FINE-LOCK approach from 1.89 to 1.87 (higher impact).

In COARSE-LOCK, we are locking entire data structure while using shared data structure. On the other side in FINE-LOCK, we are locking only specific item of the data structure while working on that piece. So as we increase the computation (Part C) at the time when we have locked the entire shared data structure, all other threads who want to work on the other part of the data structure will also wait. While in case of FINE-LOCK, two threads who want to work on different part of data structure can work simultaneously.

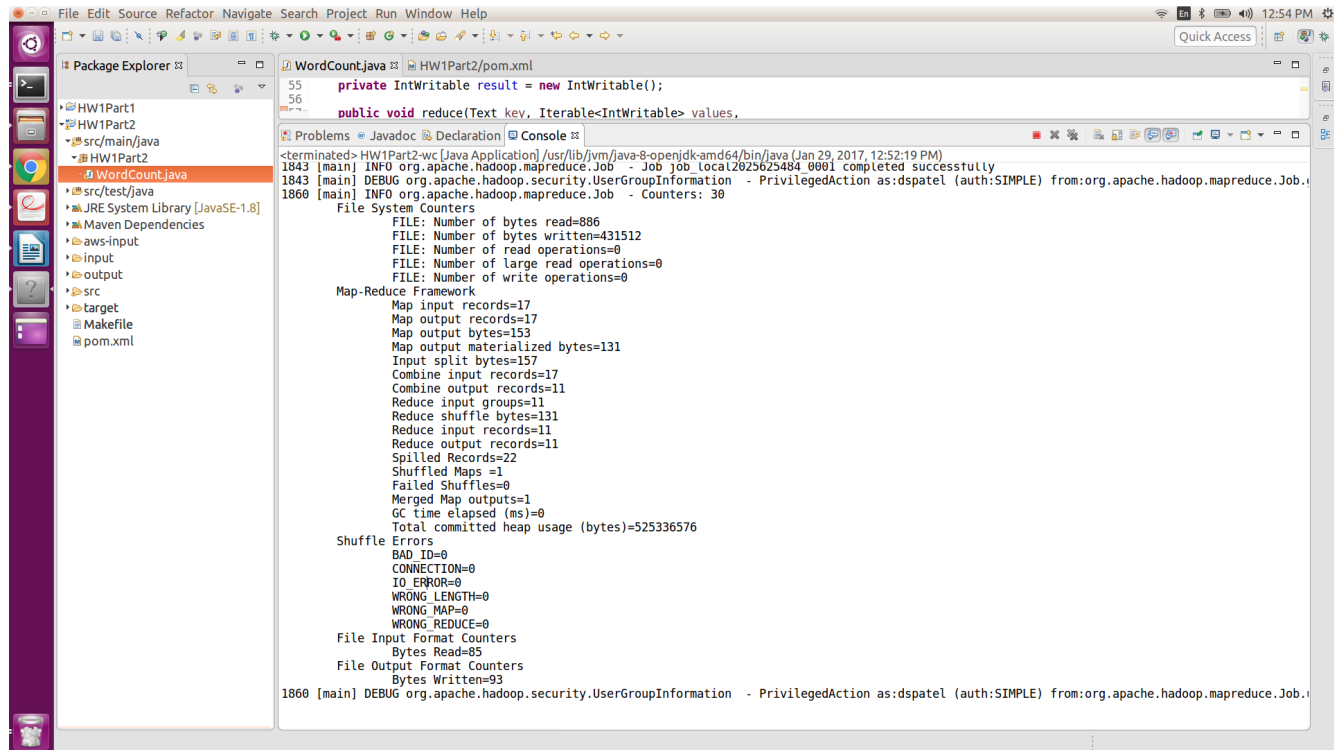
Ideally, as we increase the task to be done while locking common data structure (inside synchronized block) the performance of COARSE-LOCK approach should reduce at higher rate compared to the reduction in performance of FINE-LOCK approach.

## 2. Word Count Local Execution

### 2.1) Project directory structure



## 2.2) Console Output for a successful run of the WordCount Problem from IDE and Command Prompt



The screenshot shows an IDE window with the following components:

- Package Explorer:** Displays the project structure with folders like HW1Part1, HW1Part2, src/main/java, src/test/java, JRE System Library [JavaSE-1.8], Maven Dependencies, aws-input, input, output, src, target, Makefile, and pom.xml. The file WordCount.java is selected.
- Editor:** Shows the code for WordCount.java, including the reduce method: 

```
private IntWritable result = new IntWritable();  
public void reduce(Text key, Iterable<IntWritable> values,
```
- Console:** Displays the output of the job run, including the following information:
  - Job completion: 1843 [main] INFO org.apache.hadoop.mapreduce.Job - Job job\_local2025625484\_0001 completed successfully
  - Debug message: 1843 [main] DEBUG org.apache.hadoop.security.UserGroupInformation - PrivilegedAction as:dspatel (auth:SIMPLE) from:org.apache.hadoop.mapreduce.Job.
  - Info message: 1860 [main] INFO org.apache.hadoop.mapreduce.Job - Counters: 30
  - File System Counters:
    - FILE: Number of bytes read=886
    - FILE: Number of bytes written=431512
    - FILE: Number of read operations=0
    - FILE: Number of large read operations=0
    - FILE: Number of write operations=0
  - Map-Reduce Framework:
    - Map input records=17
    - Map output records=17
    - Map output bytes=153
    - Map output materialized bytes=131
    - Input split bytes=157
    - Combine input records=17
    - Combine output records=11
    - Reduce input groups=11
    - Reduce shuffle bytes=131
    - Reduce input records=11
    - Reduce output records=11
    - Spilled Records=22
    - Shuffled Maps =1
    - Failed Shuffles=0
    - Merged Map outputs=1
    - GC time elapsed (ms)=0
    - Total committed heap usage (bytes)=525336576
  - Shuffle Errors:
    - BAD\_ID=0
    - CONNECTION=0
    - IO\_ERROR=0
    - WRONG\_LENGTH=0
    - WRONG\_MAP=0
    - WRONG\_REDUCE=0
  - File Input Format Counters:
    - Bytes Read=85
  - File Output Format Counters:
    - Bytes Written=93
  - Debug message: 1860 [main] DEBUG org.apache.hadoop.security.UserGroupInformation - PrivilegedAction as:dspatel (auth:SIMPLE) from:org.apache.hadoop.mapreduce.Job.

```

Terminal File Edit View Search Terminal Help
dpsatel@dpsatel:~/Study/Masters/3 Spring 2017/MR/HW1/workspace/HW1Part2$ make alone
mvn clean package
[INFO] Scanning for projects...
[INFO]
[INFO] -----
[INFO] Building HW1Part2 1.0
[INFO] -----
[INFO]
[INFO] --- maven-clean-plugin:2.5:clean (default-clean) @ HW1Part2 ---
[INFO] Deleting /home/dpsatel/Study/Masters/3 Spring 2017/MR/HW1/workspace/HW1Part2/target
[INFO]
[INFO] --- maven-resources-plugin:2.6:resources (default-resources) @ HW1Part2 ---
[WARNING] Using platform encoding (UTF-8 actually) to copy filtered resources, i.e. build is platform dependent!
[INFO] skip non existing resourceDirectory /home/dpsatel/Study/Masters/3 Spring 2017/MR/HW1/workspace/HW1Part2/src/main/resources
[INFO]
[INFO] --- maven-compiler-plugin:3.3:compile (default-compile) @ HW1Part2 ---
[INFO] Changes detected - recompiling the module!
[WARNING] File encoding has not been set, using platform encoding UTF-8, i.e. build is platform dependent!
[INFO] Compiling 1 source file to /home/dpsatel/Study/Masters/3 Spring 2017/MR/HW1/workspace/HW1Part2/target/classes
[INFO]
[INFO] --- maven-resources-plugin:2.6:testResources (default-testResources) @ HW1Part2 ---
[WARNING] Using platform encoding (UTF-8 actually) to copy filtered resources, i.e. build is platform dependent!
[INFO] skip non existing resourceDirectory /home/dpsatel/Study/Masters/3 Spring 2017/MR/HW1/workspace/HW1Part2/src/test/resources
[INFO]
[INFO] --- maven-compiler-plugin:3.3:testCompile (default-testCompile) @ HW1Part2 ---
[INFO] Nothing to compile - all classes are up to date
[INFO]
[INFO] --- maven-surefire-plugin:2.17:test (default-test) @ HW1Part2 ---
[INFO] No tests to run.
[INFO]
[INFO] --- maven-jar-plugin:2.4:jar (default-jar) @ HW1Part2 ---
[INFO] Building jar: /home/dpsatel/Study/Masters/3 Spring 2017/MR/HW1/workspace/HW1Part2/target/HW1Part2-1.0.jar
[INFO]
[INFO] --- maven-shade-plugin:2.4.3:shade (default) @ HW1Part2 ---
[INFO] Replacing original artifact with shaded artifact.
[INFO] Replacing /home/dpsatel/Study/Masters/3 Spring 2017/MR/HW1/workspace/HW1Part2/target/HW1Part2-1.0.jar with /home/dpsatel/Study/Masters/3 Spring 2017/MR/HW1/work
space/HW1Part2/target/HW1Part2-1.0-shaded.jar
[INFO]
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO]
[INFO] Total time: 3.726 s
[INFO] Finished at: 2017-01-29T12:55:49-05:00
[INFO] Final Memory: 26M/200M
[INFO]
[INFO] -----
rm -rf output*
/usr/local/hadoop/bin/hadoop jar target/HW1Part2-1.0.jar HW1Part2.WordCount input output
17/01/29 12:55:52 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id

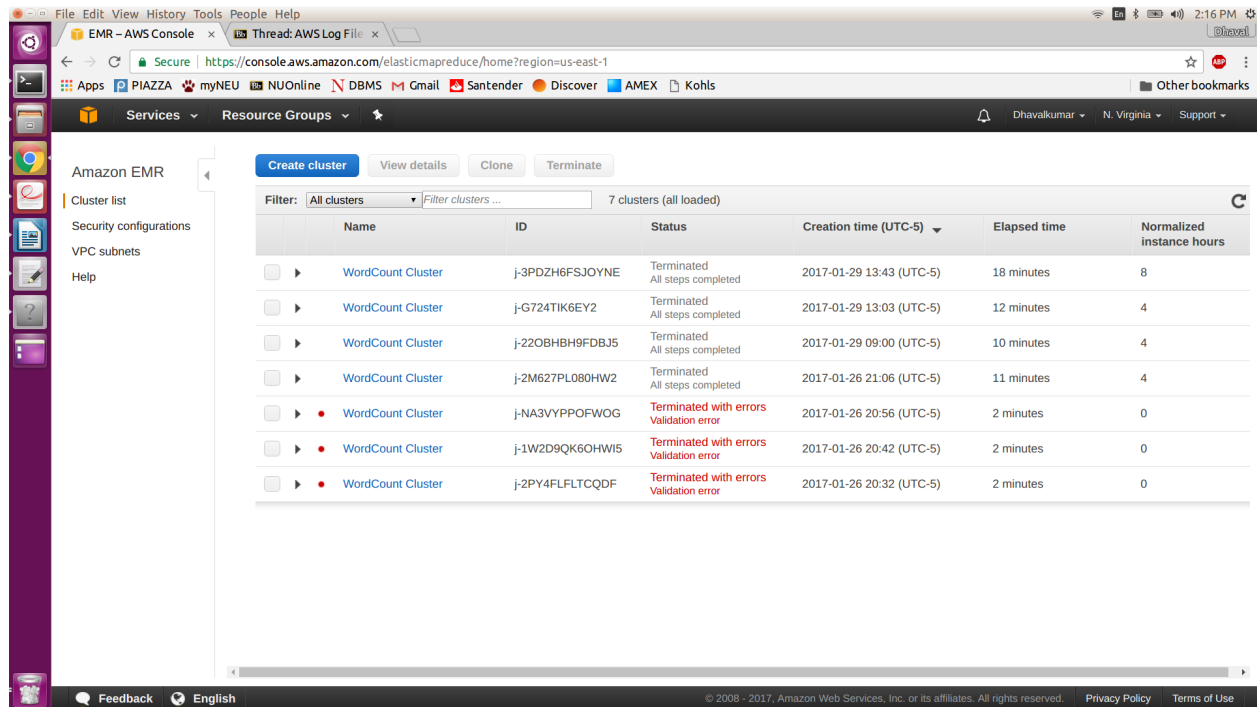
```

```

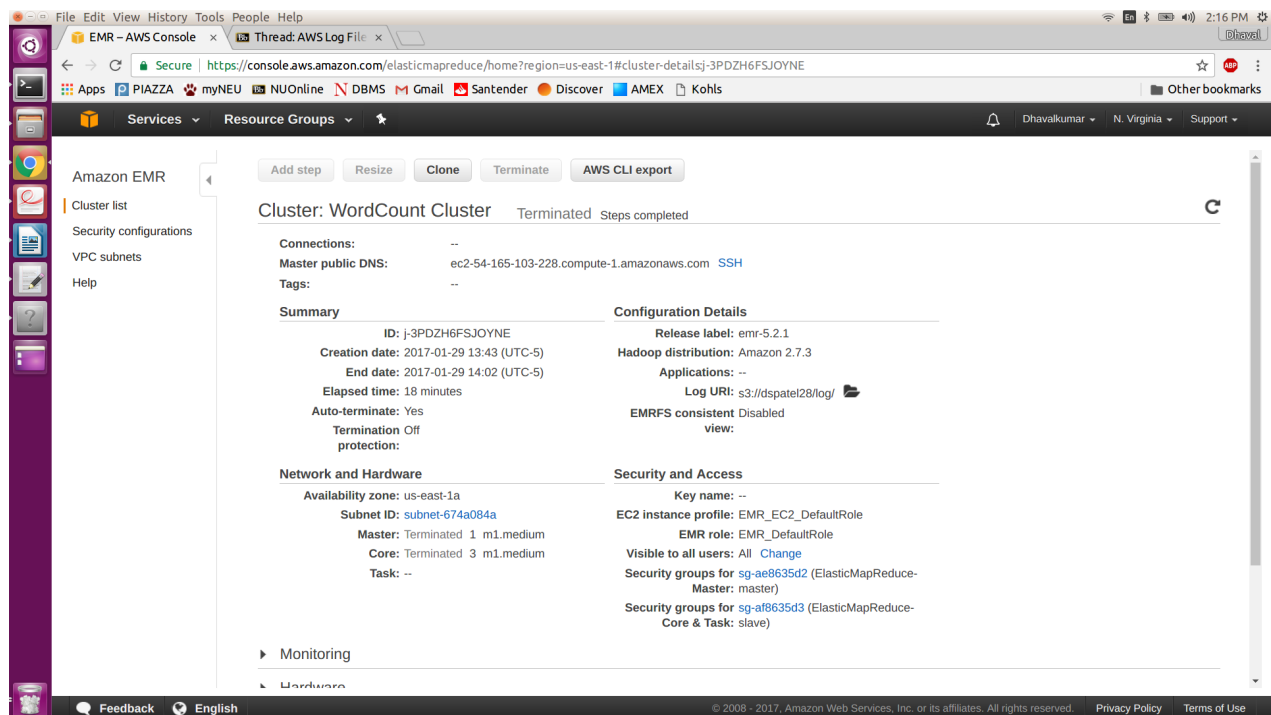
Terminal File Edit View Search Terminal Help
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=85
File Output Format Counters
  Bytes Written=93
1469 [main] INFO org.apache.hadoop.mapreduce.Job - Counters: 30
File System Counters
  FILE: Number of bytes read=11128
  FILE: Number of bytes written=582130
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
Map-Reduce Framework
  Map input records=17
  Map output records=17
  Map output bytes=153
  Map output materialized bytes=131
  Input split bytes=157
  Combine input records=17
  Combine output records=11
  Reduce input groups=11
  Reduce shuffle bytes=131
  Reduce input records=11
  Reduce output records=11
  Spilled Records=22
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=12
  Total committed heap usage (bytes)=460324864
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=85
File Output Format Counters
  Bytes Written=93
dpsatel@dpsatel:~/Study/Masters/3 Spring 2017/MR/HW1/workspace/HW1Part2$

```

### 3 Word Count AWS Execution

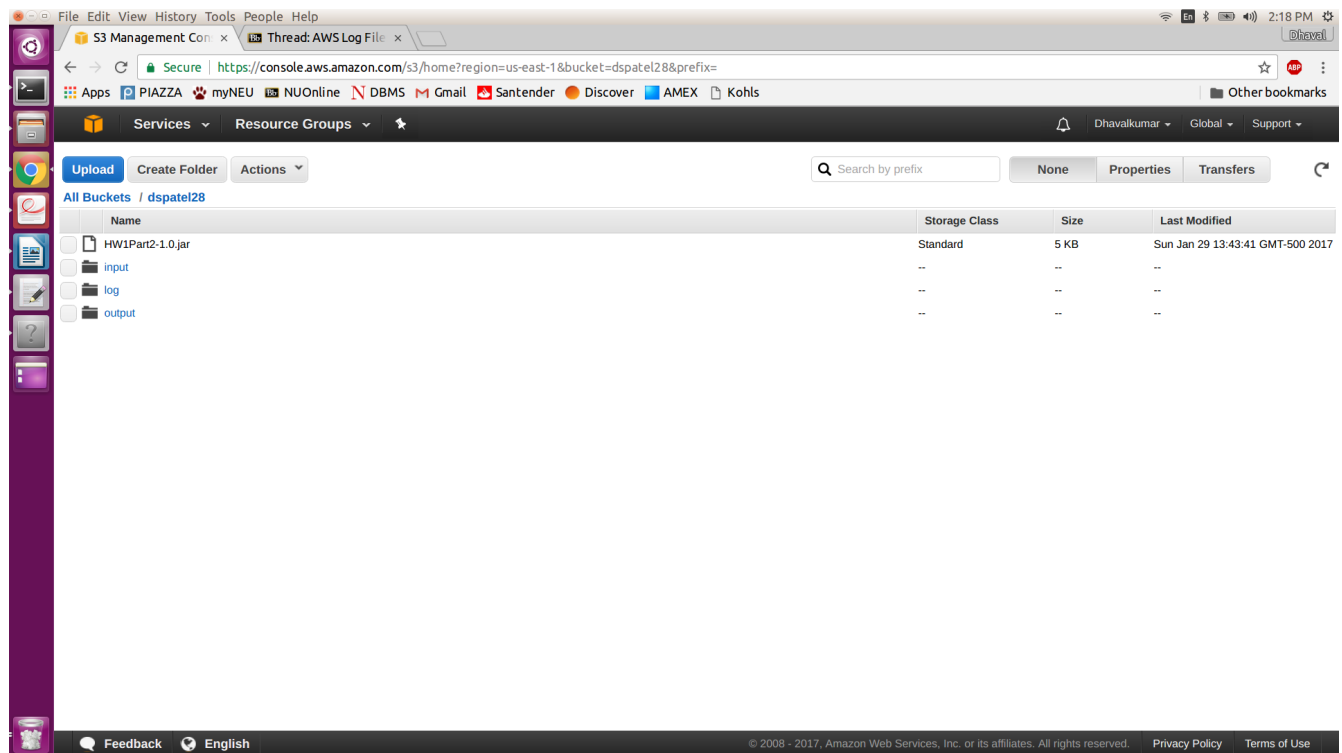


You can see 1 Master and 3 worker nodes in below image:





Bucket after successful run:



Result files after successful run:

